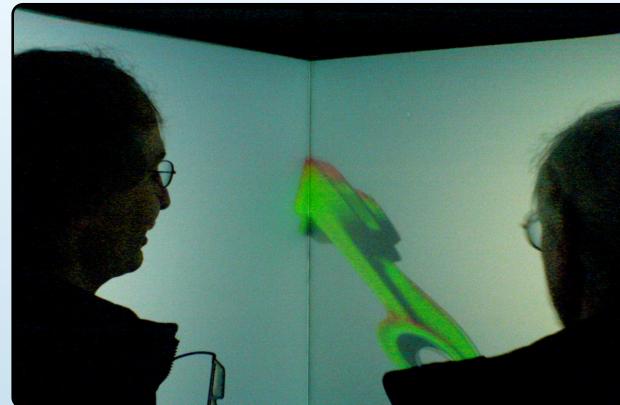


Scalability

Equalizer implements a wide range of algorithms to parallelize the rendering of large data sets. Multiple graphic cards, processors and computers can be combined to render a single view. The Equalizer server distributes the rendering task across the available resources (decomposition) and assembles the results on the final view (recomposition).

For the task decomposition, Equalizer currently supports sort-first (2D), sort-last (DB) and stereo (Eye) compounds. Time-multiplex (DPlex) is planned.

Equalizer supports virtually any recomposition algorithm, for example binary swap or direct send for sort-last, and tile gathering for sort-first rendering.



Sponsored by:



Open standard for scalable rendering
LGPL license

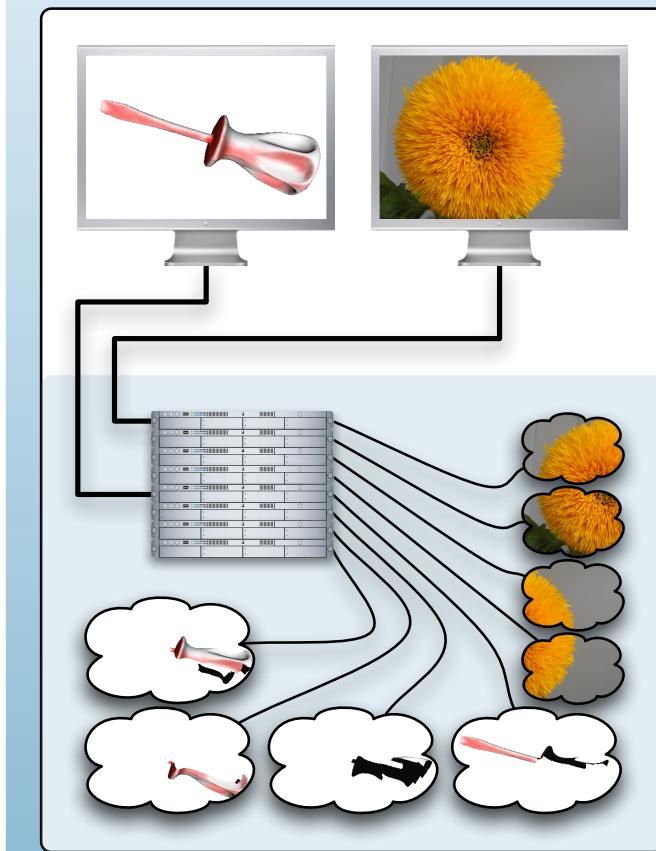
Collaborations welcome
Consulting and support available

<http://www.equalizergraphics.com>
contact@equalizergraphics.com
+41 76 33 77 247

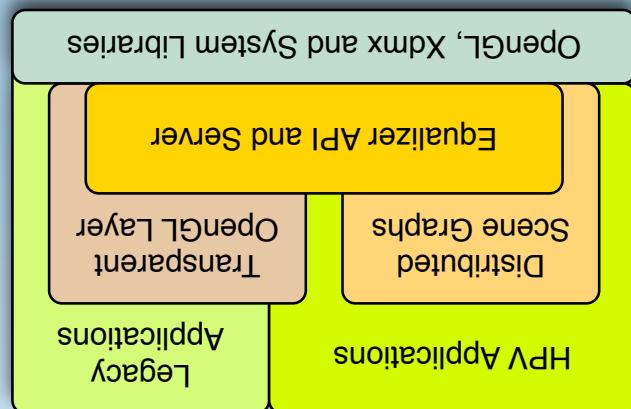
Wed Jan 03 2007
© 2006 Stefan Eilemann. All information provided is subject to change without notice.

Equalizer

Scalable Rendering



Equalizer is an open source programming interface and resource management system for scalable graphics applications. An Equalizer application can run unmodified on any visualization system, from a singlepipe workstation to large scale graphics clusters and shared memory visualization systems. The foundation of Equalizer is a parallel, scalable programming interface which solves the problems common to any multipipe application.



Resource Management

Equalizer uses a configuration server to optimally allocate and balance the available resources on the visualizaⁿtion system. The server is configured using a hierarchical structure defined configuration file to describe the available resources, and the combination of the resources for rendering. Equalizer applications can also run without a server, in which case they will automatically configure themselves for the local workstation. The central resource management enables system-wide resource reservation, allocation and scheduling, as well as the integration with other cluster software.

Transparent OpenGL Layer

A transparent OpenGL layer will enable the execution of unmodified OpenGL applications. It allows a series of unmodified OpenGL applications to run simultaneously on the same system.

Distributed Scene Graphs

Edualizer will be integrated with popular scene graphs, such as Coin3D or OpenSceneGraph. Applications using these scene graphs can easily implement parallel, scalable applications, and profit from the current and future feature set of Edualizer.

Edualizer uses a callback-driven programming model, where the application provides their rendering methods, which are called by the Edualizer framework according to the current configuration. Process and thread creation, task synchronization and network connections are externalized from the application and handled by the Edualizer framework.

Programming Interface

- **Node** - a single computer in the cluster
- **Pipe** - a graphics card and rendering thread
- **Window** - an OpenGL draw able
- **Channel** - a view port within a window
- The application subclasses these entities, and overrides methods to provide the application's rendering code. Each method corresponds to a task, and for all of them Equalizer provides a default method which implements the typical use case. This allows a quick start for the developer, which can replace the defaults gradually with his own code.

- **Channel::draw** - render using the provided frustum, viewport and range
 - **Window::init** - initialize OpenGL drawable and state
 - **Pipe::startFrame** - update frame-specific data
 - **Node::init** - initialize per node application data

Eduqalizer provides an API to build network-distributed applications. In contrast to other cluster APIs, such as MPI or PVM, this API is designed for the development of interactive applications. Certain characteristics of MPI and PVM are exploited for better performance and simpler visualization software, e.g. the frame-driven rendering, are exploited for better performance and simpler programming interface.

- Multipe Workstations are an affordable way to scale the rendering performance and display size.
- Remote Visualizatoin Clusters are a cost-efficient way of centralizing large scale data.
- Scalable Rendering is used to parallelize the rendering of a single view across multiple graphic cards and processors.

Use Cases

