

# Equalizer

VizSIG Meeting, October 2006  
Stefan Eilemann

# Outline

---

- High performance visualisation (HPV)
- Equalizer
  - Programming interface
  - Resource management
  - Future components

# HPV

---

- Transparent and semi-transparent solutions
- Programming interfaces
  - Scene graphs
  - Generic middleware

# HPV Transparent Solutions

---

- Chromium, ModViz VGP, OMP
- Operate on OpenGL command stream
- Programming extensions for improved performance and scalability (semi-transparent)
- HPC analogy: auto-parallelising compilers

# Scene Graph API's

---

- ScaleViz, Vega Prime, VTK, OpenSG
- Impose overall programming model and data structure
- Best for developing new applications
- HPC analogy: CFD codes

# Generic HPV Middleware

---

- Cavelib, VRJuggler, MPK
- Limited to HPV-critical areas of the code
- Best for porting existing applications
- HPC analogy: MPI, PVM

# Equalizer

---

A Programming Interface

*and*

Resource Management System

*for*

Scalable Graphics Applications

# Equalizer Programming Interface

---

Applications are written against a *client library* which abstracts the interface to the execution environment

- Minimally invasive programming approach
- Abstracts multi-processing, synchronisation and data transport
- Supports distributed rendering and performs frame compositing

# Equalizer Programming Interface

---

C++ classes which correspond to graphic entities,  
e.g.:

- **Node** – a single computer in the cluster
- **Pipe** – a graphics card and rendering *thread*
- **Window** – an OpenGL drawable
- **Channel** – a viewport within a window

# Equalizer Programming Interface

---

Application subclasses and overrides methods,  
e.g.:

- **Channel::draw** – render using the provided frustum, viewport and range
- **Window::init** – init OpenGL drawable and state
- **Pipe::startFrame** – update frame-specific data
- **Node::init** – initialise per node application data

Default methods implement typical use case

# Resource Management System

---

Applications are deployed by a *server* which balances the resource usage across the system

- Centralises the setup for all applications
- Configures application and deploys render clients
- Dynamic load-balancing of the cluster resources

# Resource Management System

---

## Server configuration:

```
server
{
    config // 1-n times, currently only the first one is used by the server
    {
        <resources> // What is being used (next slide)
        <compounds> // How it is being used (slide after next)
    }
}
```

# Resource Management System

---

## Resource configuration:

```
node // 1-n times
{
    pipe // 1-n times
    {
        display  unsigned      // X11 display or ignored
        screen   unsigned      // X11 screen/CGL display/graphics adapter
        window  // 1-n times
        {
            viewport [ viewport ] // wrt pipe, default full screen
            channel // 1-n times
            {
                name    string
                viewport [ viewport ] // wrt window, default full window
            }
        }
    }
}
```

# Resource Management System

---

## Resource utilisation:

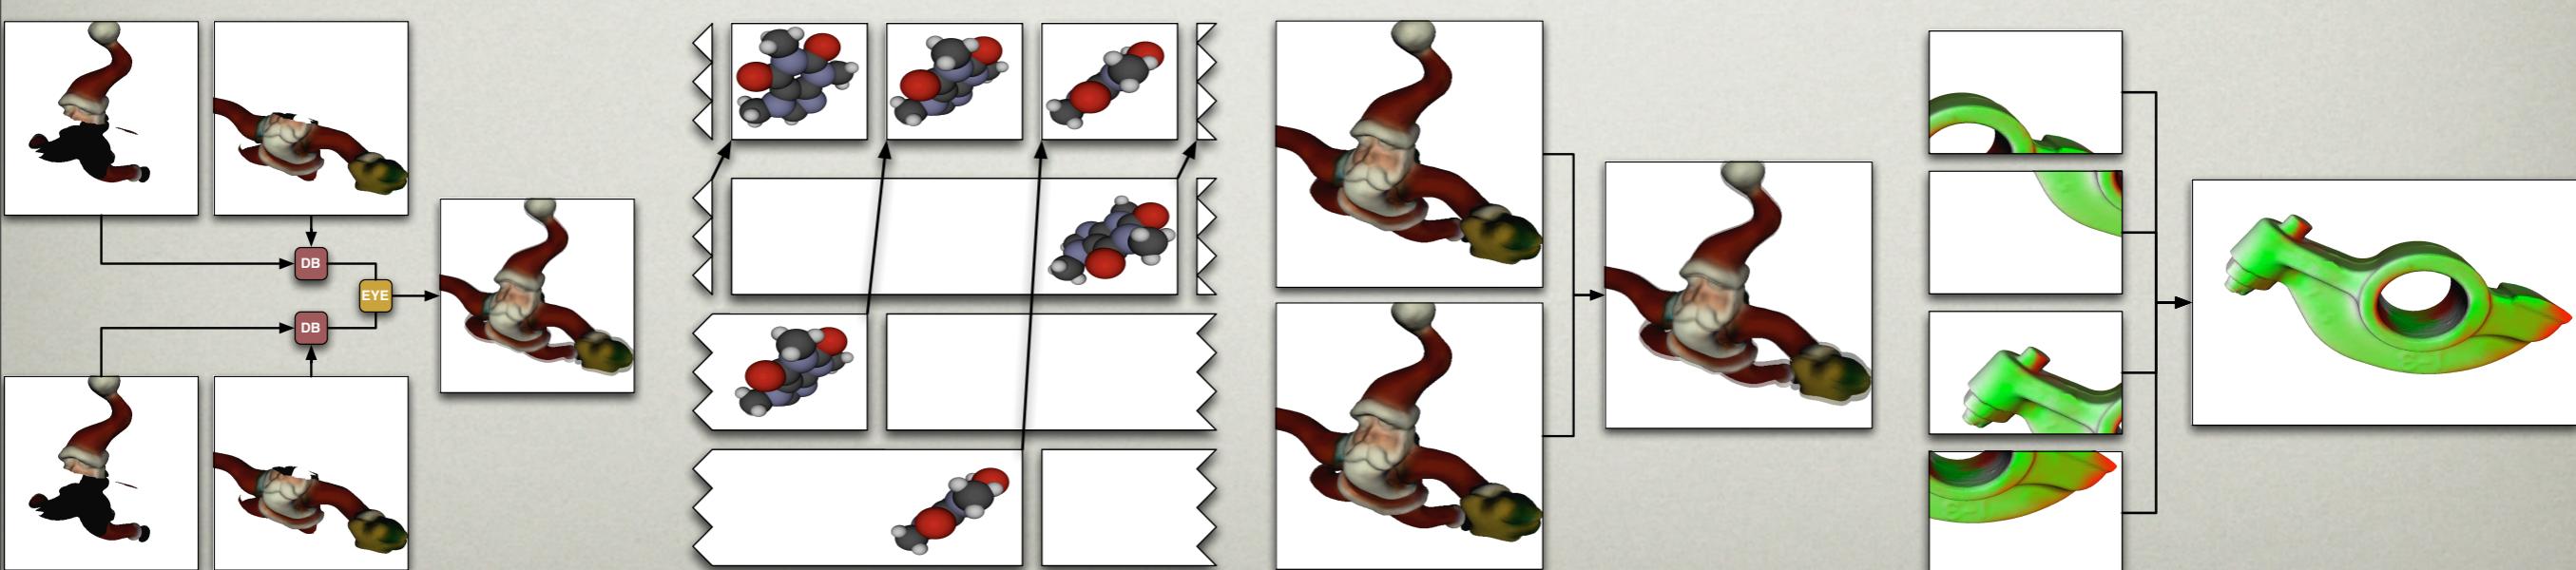
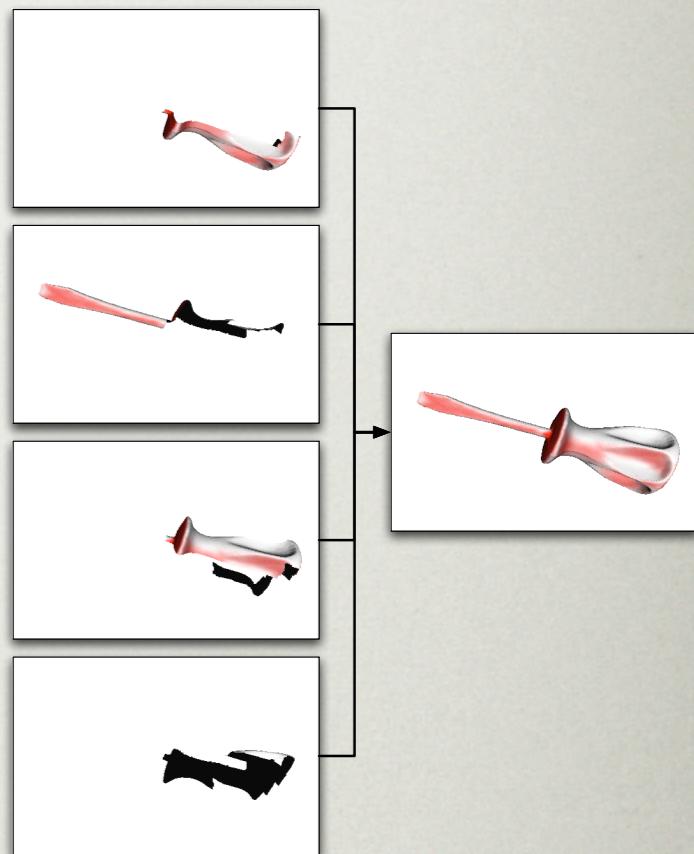
```
compound // 1-n times
{
    channel string    // where the compound tasks are executed
    task      [ CLEAR DRAW ASSEMBLE READBACK ] // tasks to execute
    viewport [ viewport ]           // wrt parent compound, sort-first
    range     [ float float ]        // DB-range for sort-last
    eye       [ CYCLOP LEFT RIGHT ]   // monoscopic or stereo view
    wall|projection          // frustum description
    {...}                      // typically at root compound

    <child-compounds>

    swapBarrier { name string } // same barriername = sync swap buffers
    outputFrame { name string }
    inputFrame  { name string } // name corresponding to an output frame
}
```

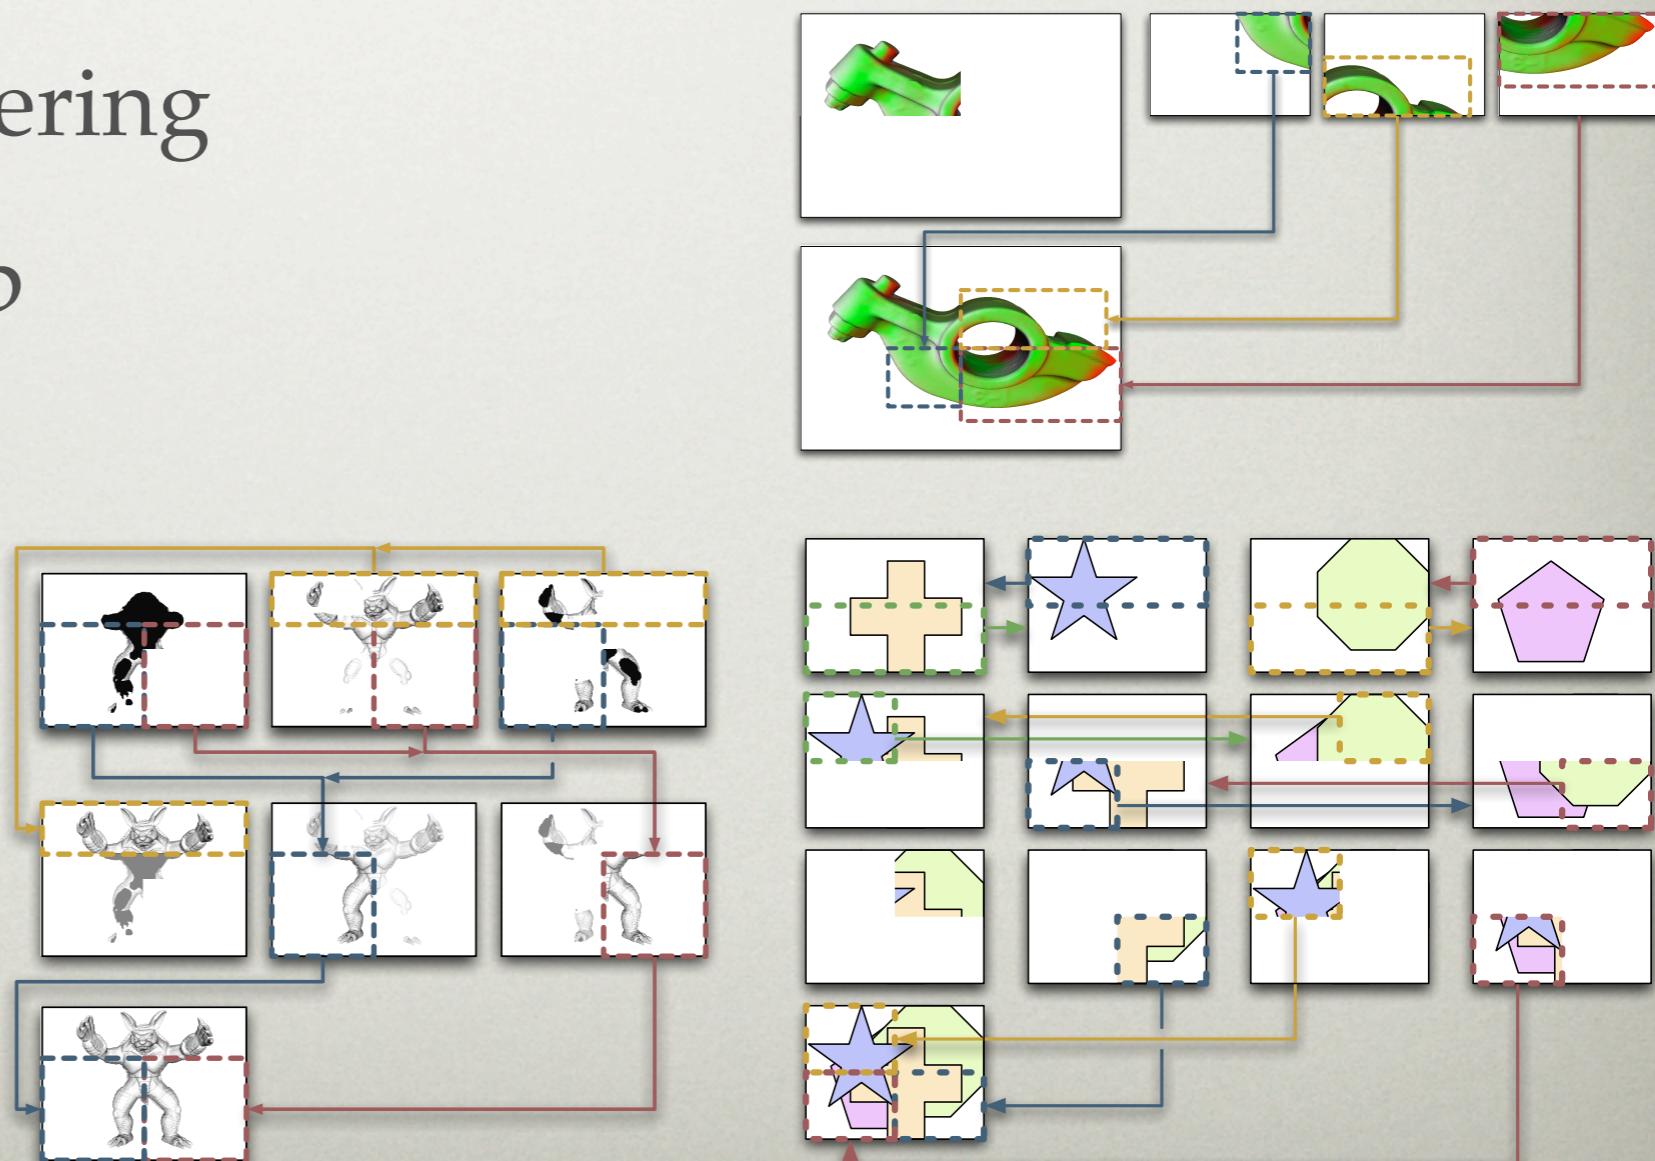
# Decomposition Modes

- DB / sort-last (range)
- 2D / sort-first (viewport)
- Eye / stereo (eye)
- DPlex (period, phase)
- Any combination thereof



# Recomposition Modes

- Combination of task and frames allows virtually any recomposition mode, e.g.:
  - 2D tile gathering
  - binary swap
  - direct send

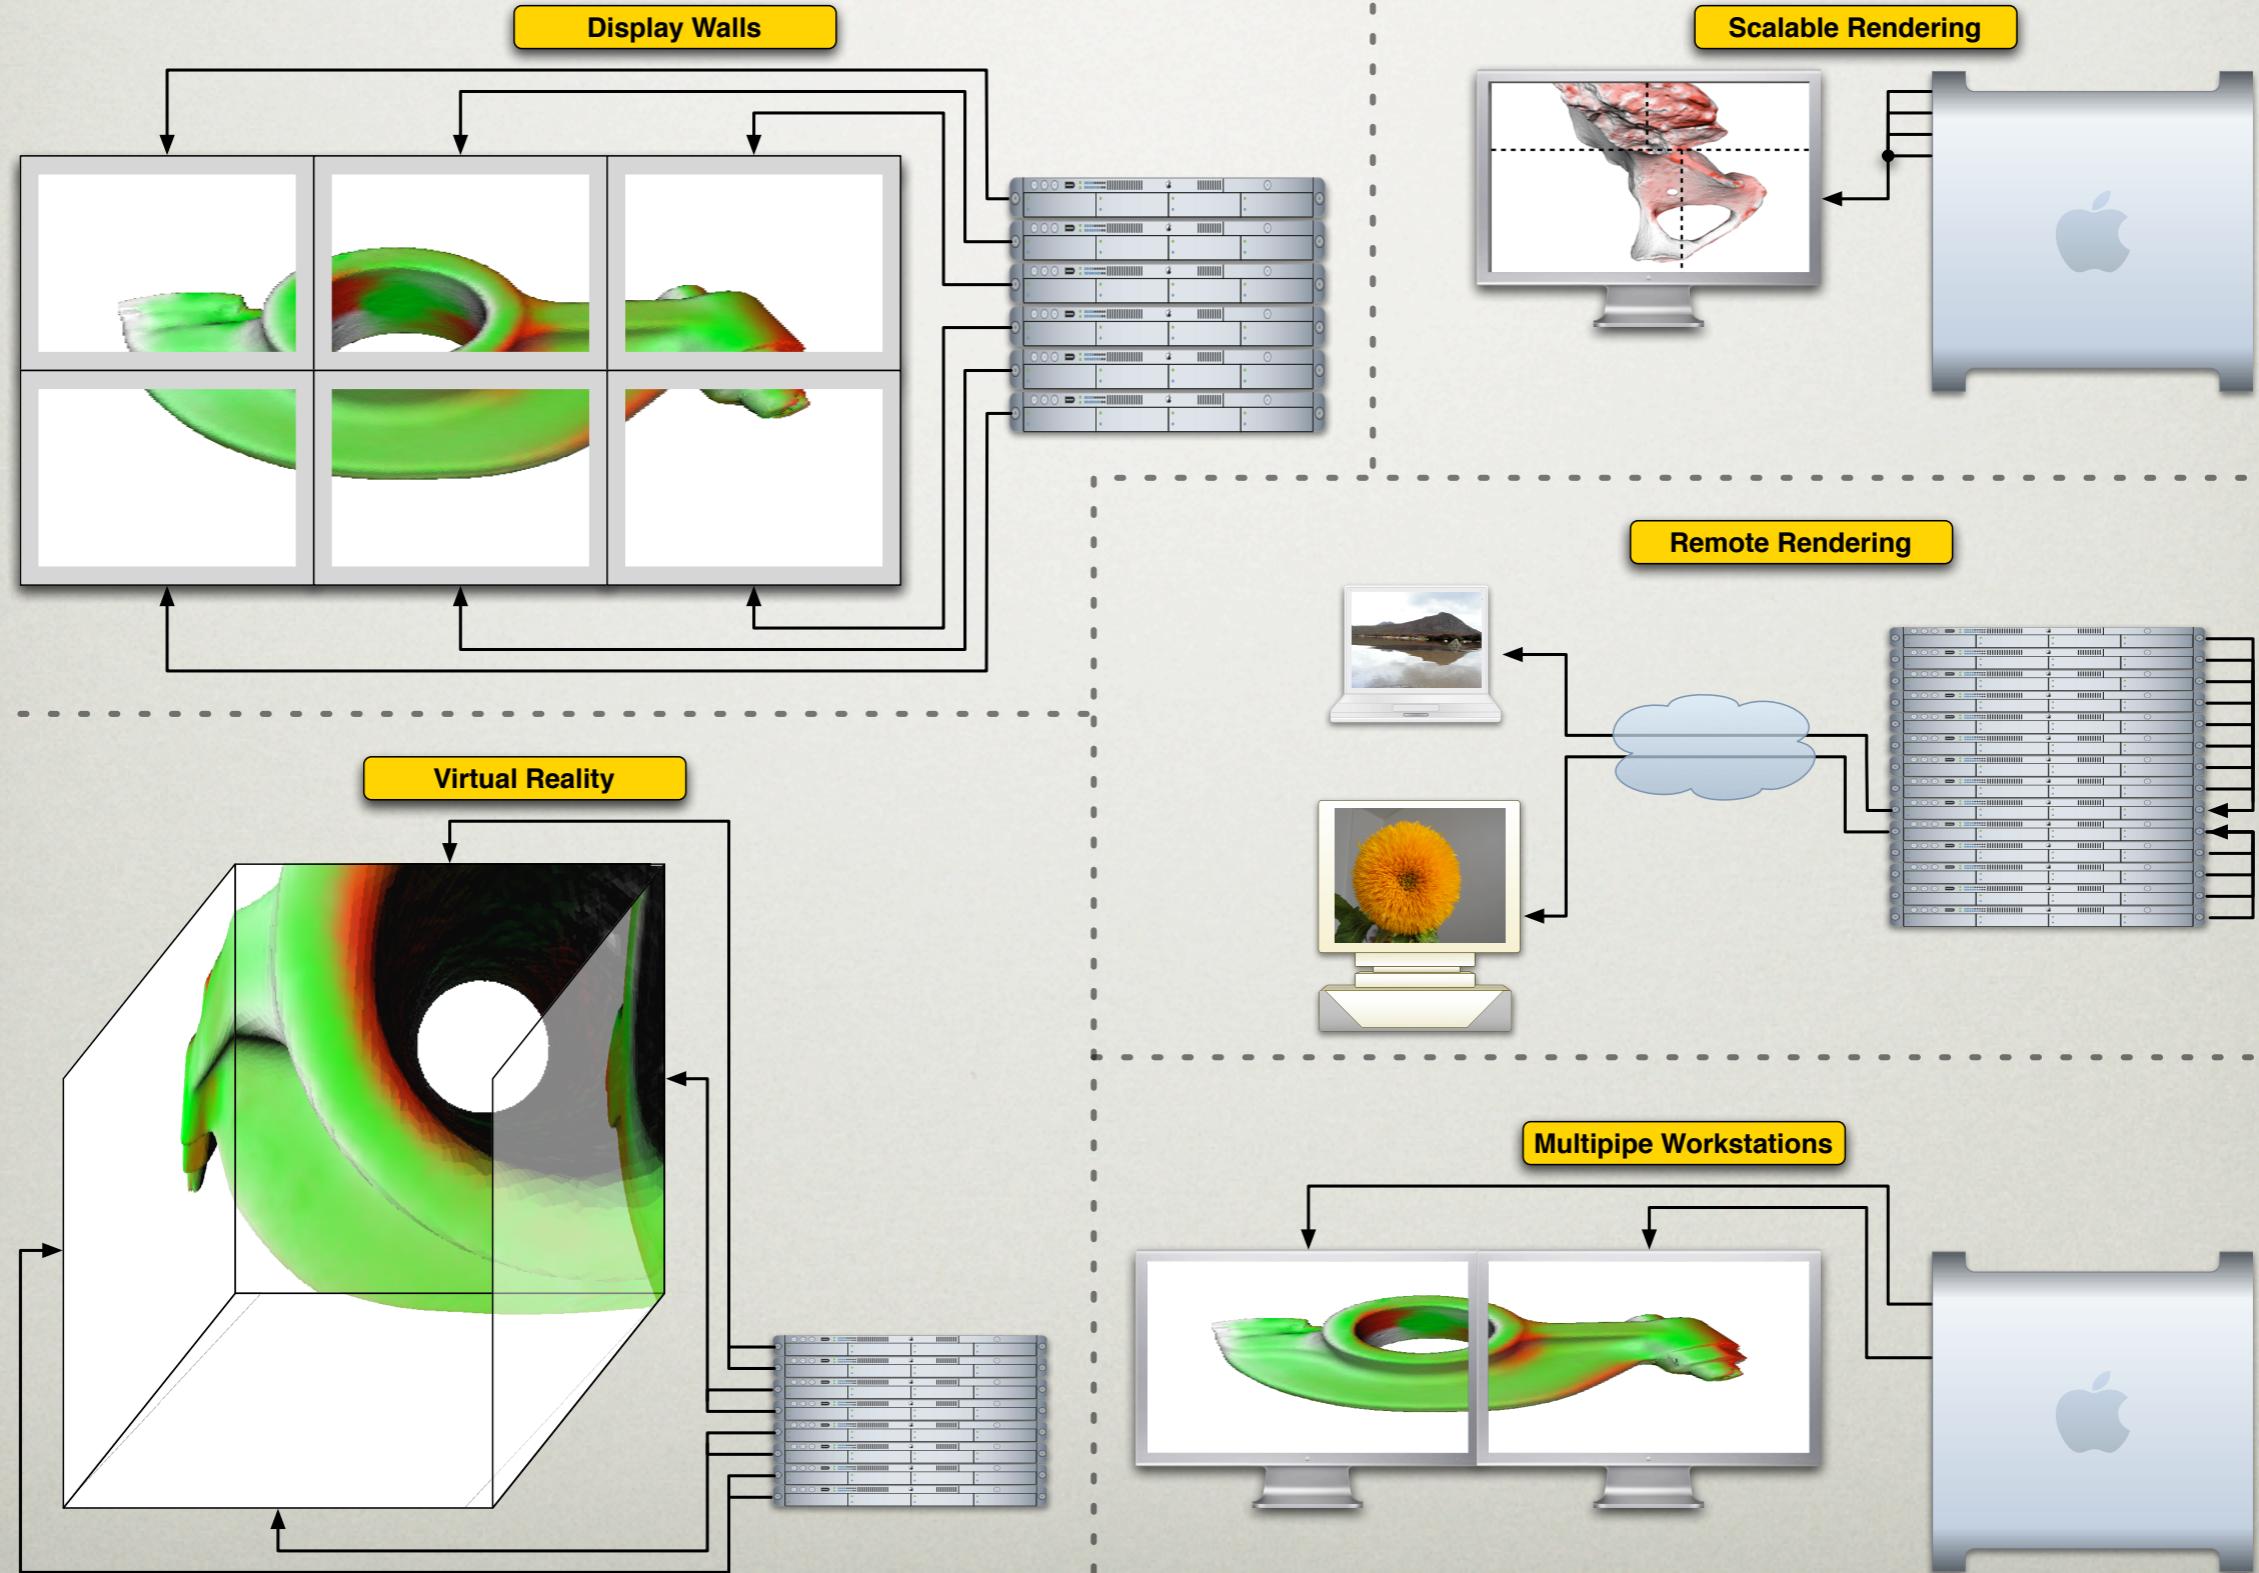


# SSI and Clusters

---

- Supercomputers are just tightly integrated clusters
- Equalizer runs on both architectures
- Execution model is the same
- SSI allows additional optimisations and simplifications
- Stand-alone SSI version planned

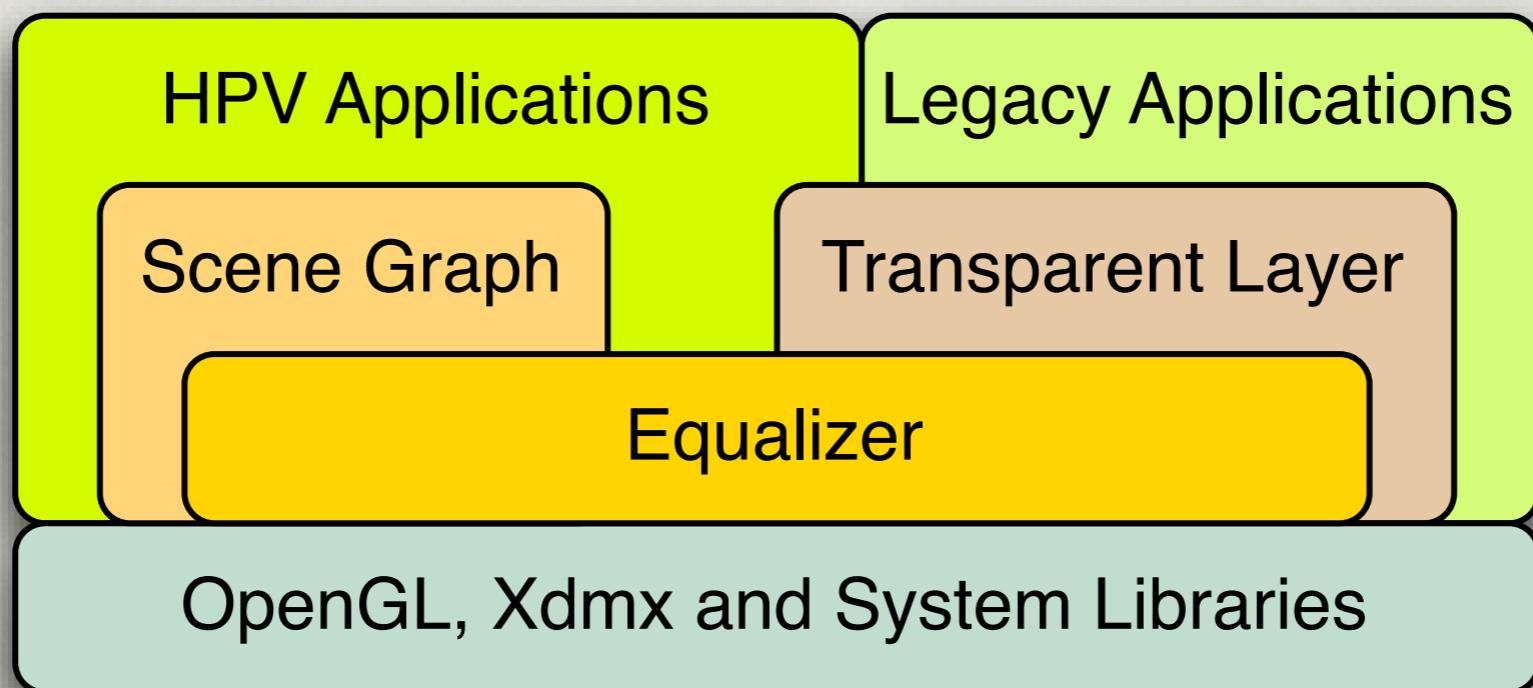
# Use Cases



# Equalizer Future

---

- Transparent Layer: virtual OpenGL screen
- Scene Graph: “transparent scalability”
- Equalizer: Scalable rendering engine



# Transparent Layer

---

- “Enabler” for visualisation clusters
- Legacy and non-critical applications
- System load balancing between transparent and HPV applications
- Single point of configuration
- Performance and compatibility as today

# Distributed Scene Graph

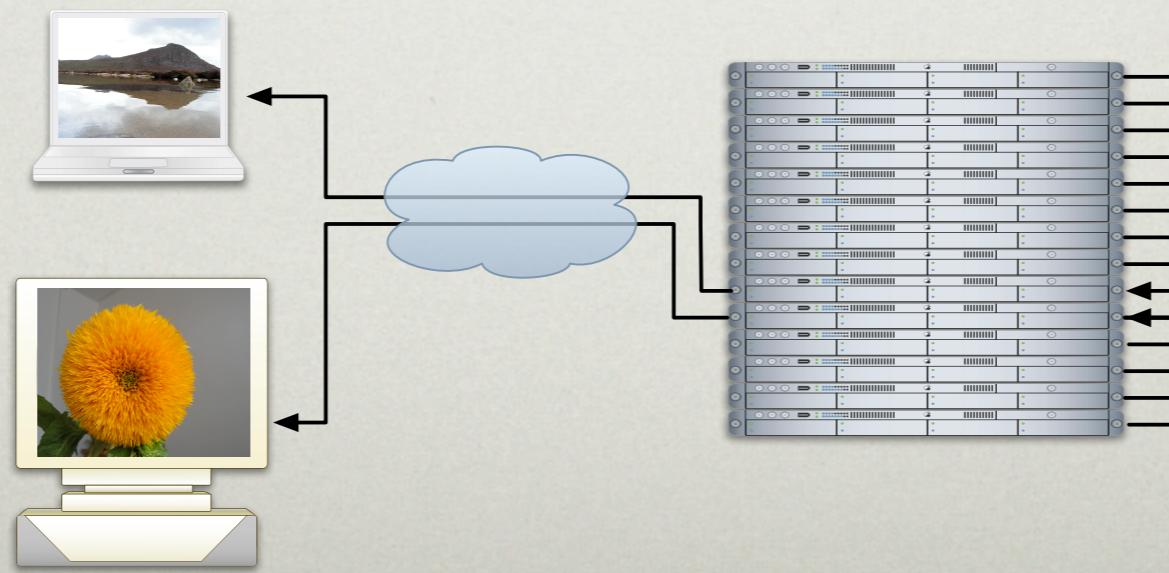
---

- Uses Equalizer for HPV
- Small effort for application developer
- Single point of configuration
- Candidates: OpenSceneGraph, Coin

# Remote Visualisation

---

- Leverages knowledge of the application
  - Frames are often available in main memory
  - Additional frame-transport optimisations
- Loadbalancing of multiple applications on one visualisation cluster



# Project Status

---

- API and resource server are usable
- Transparent layer, remote visualisation and distributed scene graph depend on demand and sponsoring

# Last Words

---

- LGPL license
- Open standard for scalable graphics
- User-driven development
- Alpha version available on:  
[www.equalizergraphics.com](http://www.equalizergraphics.com)
- Consulting and support available
- [Get in touch](#)

# Demos

---

- 1-window
- 3-pipe
- 5-channel.cave